

# CS542200 Parallel Programming

## HW4: Blocked All-Pairs Shortest Paths

Josh Kao

National Tsing Hua University

2014/12/11

# Outline I

- 1 Problem Description
- 2 Input/Output Formats
- 3 Working items
- 4 Grading
- 5 Reminder

# All-Pairs Shortest Paths

Given an  $N \times N$  matrix  $D = [d(i, j)]$  where  $d(i, j)$  represents the shortest-path distance from a vertex  $i$  to a vertex  $j$ .

Let  $D^{(k)} = [d^{(k)}(i, j)]$  be the result which all the intermediate vertices are in the set  $\{1, 2, \dots, k\}$ .

## Floyd-Warshall Method

$$d^{(k)}(i, j) = \begin{cases} \text{weight of } (i, j) & \text{if } k = 0; \\ \min(d^{(k-1)}(i, j), d^{(k-1)}(i, k) + d^{(k-1)}(k, j)) & \text{if } k \geq 1. \end{cases}$$

The matrix  $D^{(N)} = [d^{(N)}(i, j)] = D$  gives an answer to APSP problem.

# FW Algorithm

## Pseudo Code

```
for  $k \leftarrow 1$  to  $N$  do  
  | for  $i \leftarrow 1$  to  $N$  do  
    | for  $j \leftarrow 1$  to  $N$  do  
      |  $d[i][j] \leftarrow \min(d[i][j], d[i][k] + d[k][j]);$   
    | end  
  | end  
end
```

# Parallel Strategy

- For each  $\mathbf{k}$ , calculate  $d^{(k)}(i, j)$  in parallel

## FW algorithm CUDA version

```
for  $k \leftarrow 1$  to  $N$  do  
  | FW_kernel<<< #blocks, #threads >>>( ... );  
end
```

# But ....

Q. How can we do APSP on multi-GPUs **individually**?

# Blocked FW algorithm

- Partition matrix into  $N/B \times N/B$  blocks of  $B \times B$  submatrices.
- For intance,  $\mathbf{N} = 6$ ,  $\mathbf{B}$  (Blocking Factor) = 2 :

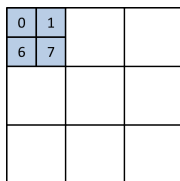
0	1	2	3	4	5
6	7				

Block (1,1)	Block (1,2)	Block (1,3)
Block (2,1)	Block (2,2)	Block (2,3)
Block (3,1)	Block (3,2)	Block (3,3)

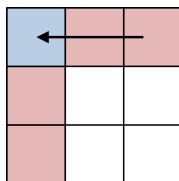
Figure: Divide a matrix by  $B = 2$

# Blocked FW algorithm

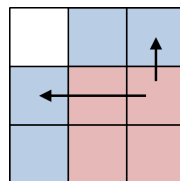
- 3 phases in each iteration
  - Phase 1: Self-dependent block
  - Phase 2: Pivot-row & Pivot-column blocks
  - Phase 3: Other blocks



(a) Phase 1



(b) Phase 2



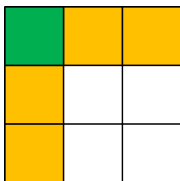
(c) Phase 3

Figure: Blocked algorithm in 1<sup>st</sup> iteration

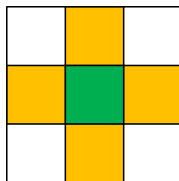


# Blocked FW algorithm

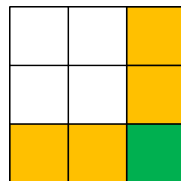
- It will run  $N/B$  iterations
- In this example, it will run  $6/2 = 3$  iterations.
  - Phase 1: **Green**; Phase 2: **Yellow**; Phase 3: **White**



(a) Iteration 1



(b) Iteration 2



(c) Iteration 3

Figure: Blocked algorithm

# Outline I

- 1 Problem Description
- 2 Input/Output Formats**
- 3 Working items
- 4 Grading
- 5 Reminder

# Execution Format

- Your program has to be able to **reading file**, and generate output in another file.
- Your program accepts two input parameters.

## Format

```
$ ./{YOUR PROGRAM} {INPUT FILE} {OUTPUT FILE}
```

For instance,

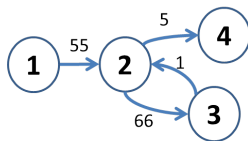
```
[kjs1095@lsalab ~]$ ./HW4_103065566_cuda.exe tiny_testcase output  
[kjs1095@lsalab ~]$ diff -b output answer
```

# Input Format

- The test case begins with a line contains 2 integers **N** and **M**
  - N : num of vertices ( $1 \leq N \leq 6000$ )
  - M : num of **directed** edges
- Each of the following **M** lines contains 3 integers **S**, **T** and **D**
  - S: source vertex id ( $S \neq T$ )
  - T: target vertex id
  - D: distance from **S** to **V** ( $0 \leq D \leq 100$ )

```
[kjs1095@lsalab ~]$ cat tiny_testcase
4 4
1 2 55
2 3 66
3 2 1
2 4 5
```

(a) Sample content of input



(b) Graph

Figure: Sample Input

# Output Format

- Show the shortest distance from  $i^{th}$  vertex to  $j^{th}$  vertex
- If there is no path from  $i^{th}$  vertex to  $j^{th}$  vertex, output **INF**

```
[kjs1095@lsalab ~]$ cat output
0 55 121 57
INF 0 66 5
INF 1 0 71
INF INF INF 0
```

Figure: Sample Output

# Outline I

- 1 Problem Description
- 2 Input/Output Formats
- 3 Working items**
- 4 Grading
- 5 Reminder

# Implementation

- Single GPU
- Multi GPU (MPI version)
- Multi GPU (OpenMP version)
- Makefile

# Report

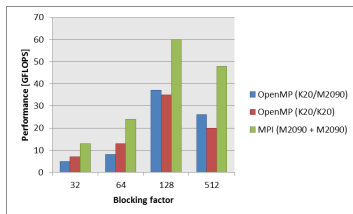
- **Implementation**
- **Profiling results**
- **Experiments**
  - ① Weak Scalability & Time Distribution
  - ② Blocking Factor
  - ③ Compare three implementations
  - ④ Others

Think of the trend of the figure and **explain**.

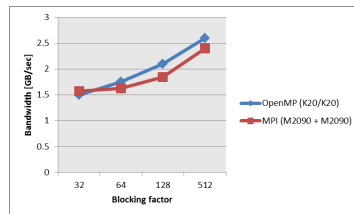
- **Experience/Conclusion**



# Example figures



(a) Performance trend



(b) Bandwidth trend

**Figure:** Example figures of blocking factor experiment

# Outline I

- 1 Problem Description
- 2 Input/Output Formats
- 3 Working items
- 4 Grading**
- 5 Reminder

# Grading

- ① [ 4%] Lab2
- ② [48%] Correctness
- ③ [30%] Report
- ④ [20%] Demo
- ⑤ [ 5%] Bonus

**Total score** = *min*( (1)+(2)+(3)+(4)+(5) , 100)

For each detail, check out the document on iLMS.

# Outline I

- 1 Problem Description
- 2 Input/Output Formats
- 3 Working items
- 4 Grading
- 5 Reminder

# Login to server

- **Host:** 140.114.91.176
- **Account:** Please refer to Account\_Table.pdf
- **Password:** 1234userX [X=1 ~ 67] (Default)

It's recommended to use ***passwd*** to change your password.

# How to run MPI version?

- Execute **setting.sh** for connecting other machines
  - `$ . setting.sh`
- Modify **hostfile**
  - Choose machines from `gpucluster0 ~ gpucluster3`
- Modify **Makefile**
  - I used `HW4_cuda_mpi.cu` for example
  - You have to change it to your own file name.

Compile MPI version (Take `HW4_cuda_mpi.cu` for example)

```
$ make HW4_cuda_mpi.exe
```

Run MPI version

```
$ mpirun -np 2 -hostfile hostfile ./HW4_cuda_mpi.exe in1 out1
```

# Reminder

- Upload **HW4\_{Student-ID}.zip** to iLMS before **1/11(Sun) 23:59:59**
  - ① HW4\_{Student-ID}\_cuda.cu
  - ② HW4\_{Student-ID}\_mpi.cu
  - ③ HW4\_{Student-ID}\_openmp.cu
  - ④ HW4\_{Student-ID}\_report.pdf
  - ⑤ Makefile
- Please **start your work ASAP** and do not leave it until the last day!
- Late submission penalty policy please refer to syllabus.
- Asking questions on iLMS or through e-mail is also welcome!